

# Homework\_1

July 16, 2021

## 0.1 MSDS 593 EDA and Visualization: Homework #1

By: Amanda Li Luo

Student ID #: 20645578

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

%config InlineBackend.figure_format = 'retina'
```

### 0.1.1 Question 1: Create a 1D ndarray of numbers from 100 to 112 (step=1) inclusively and use python to

```
[2]: x = np.arange(100, 113, 1)
x
```

```
[2]: array([100, 101, 102, 103, 104, 105, 106, 107, 108, 109, 110, 111, 112])
```

(a) Print the shape of the array

```
[3]: print('Shape:', x.shape)
```

Shape: (13,)

(b) print the data type of the array

```
[4]: print('Data type:', type(x))
```

Data type: <class 'numpy.ndarray'>

(c) Create a new array that is a slice of the original array with index [5 : 10] inclusively, and assign all values of the new array to be 0.

```
[5]: new_x = x[5:11]
print(new_x)
new_x[:] = 0
new_x
```

```
[105 106 107 108 109 110]
```

```
[5]: array([0, 0, 0, 0, 0, 0])
```

**(d) Create a boolean vector from the array to indicate if any element is greater than 105, and less than or equal to 110.**

```
[6]: boolean_x = (x > 105) & (x <= 110)

boolean_x
```

```
[6]: array([False, False, False, False, False, False, False, False, False,
        False, False, False, False])
```

**(e) Replace all the elements in the array that are greater than 105 and less than or equal to 110 with 0.**

```
[7]: x[boolean_x] = 0
```

```
[8]: x
```

```
[8]: array([100, 101, 102, 103, 104,  0,  0,  0,  0,  0,  0, 111, 112])
```

**0.1.2 Question 2: Make a Series object with year values: 1991, 1992, 1993, 1994, 1995, 1996, 1997, 1998, 1999, 2000.**

```
[9]: year = pd.Series(range(1991, 2001, 1))
year
```

```
[9]: 0    1991
     1    1992
     2    1993
     3    1994
     4    1995
     5    1996
     6    1997
     7    1998
     8    1999
     9    2000
dtype: int64
```

**(a) Print out how many total values there are using code not manual counting**

```
[10]: print('Total number of years:', year.count())
```

Total number of years: 10

**(b) Make another series with rainfall values 12.09, 12.35, 12.51, 10.25, 10.18, 10.59, 10.26, 10.48, 8.67, 10.23**

```
[11]: rainfall = pd.Series([12.09, 12.35, 12.51, 10.25, 10.18, 10.59, 10.26, 10.48, 8.67, 10.23])
rainfall
```

```
[11]: 0    12.09
      1    12.35
      2    12.51
      3    10.25
      4    10.18
      5    10.59
      6    10.26
      7    10.48
      8     8.67
      9    10.23
      dtype: float64
```

**(c) Imagine the order of rainfall values follow the order the years. Print out the years for which rainfall was less than 11.**

```
[12]: # a= rainfall[rainfall < 11].index
      # print(a)
      # year[a]
      print(year[rainfall[rainfall < 11].index].values)
```

```
[1994 1995 1996 1997 1998 1999 2000]
```

**(d) Normalize that rainfall series by xmean std and print it out.**

```
[13]: rainfall.describe()
```

```
[13]: count    10.000000
      mean     10.761000
      std      1.200106
      min      8.670000
      25%     10.235000
      50%     10.370000
      75%     11.715000
      max     12.510000
      dtype: float64
```

```
[14]: df_mean = rainfall.describe()['mean']
      df_std = rainfall.describe()['std']

      df_mean, df_std
```

```
[14]: (10.761000000000001, 1.2001060138356296)
```

```
[15]: rainfall.index
```

```
[15]: RangeIndex(start=0, stop=10, step=1)
```

```
[16]: for num in range(0,10):
        rainfall.iloc[num] = (rainfall.iloc[num] - df_mean)/ df_std

print(rainfall)
```

```
0    1.107402
1    1.324050
2    1.457371
3   -0.425796
4   -0.484124
5   -0.142487
6   -0.417463
7   -0.234146
8   -1.742346
9   -0.442461
dtype: float64
```

**(e) Set the year starting 1996 and forward as np.nan. Count the number of missing values.**

```
[17]: year[year >= 1996] = np.nan
missing_num = year.isna().value_counts()[False]
print('Number of missing values:', missing_num)
```

Number of missing values: 5

**(f) Fill all the NaNs with 0.**

```
[18]: year[year.isna()] = 0
year
```

```
[18]: 0    1991.0
1    1992.0
2    1993.0
3    1994.0
4    1995.0
5         0.0
6         0.0
7         0.0
8         0.0
9         0.0
dtype: float64
```

### 0.1.3 Question 3: For the cars.csv data set:

```
[19]: cars = pd.read_csv('cars.csv')
cars.head(10)
```

```
[19]:
```

	MPG	CYL	ENG	WGT
0	18.0	8	307.0	3504
1	15.0	8	350.0	3693
2	18.0	8	318.0	3436
3	16.0	8	304.0	3433
4	17.0	8	302.0	3449
5	15.0	8	429.0	4341
6	14.0	8	454.0	4354
7	14.0	8	440.0	4312
8	14.0	8	455.0	4425
9	15.0	8	390.0	3850

(a) Use numpy's function `np.corrcoef(x, y)` to compute the correlation between a car's weight and the miles per gallon; that function returns a matrix of  $x$  with  $x$ ,  $x$  with  $y$ , etc... so the diagonal will always be correlation 1.0. What is the correlation between a car's weight and the MPG? What does the correlation tell us about their relationship?

The correlation between a car's weight and the miles per gallon is -0.83224421. This correlation is a strong negative correlation indicating that as a car's weight increases, the number of miles per gallon decreases (and vice versa).

```
[20]: np.corrcoef(cars['MPG'], cars['WGT'])
```

```
[20]: array([[ 1.          , -0.83224421],
          [-0.83224421,  1.          ]])
```

(b) Display the records for all 8 cylinder cars

```
[21]: cars[cars['CYL'] == 8]
```

```
[21]:
```

	MPG	CYL	ENG	WGT
0	18.000000	8	307.0	3504
1	15.000000	8	350.0	3693
2	18.000000	8	318.0	3436
3	16.000000	8	304.0	3433
4	17.000000	8	302.0	3449
...	...	...	...	...
289	19.200001	8	267.0	3605
290	18.500000	8	360.0	3940
296	23.000000	8	350.0	3900
298	23.900000	8	260.0	3420
359	26.600000	8	350.0	3725

[103 rows x 4 columns]

(c) Create a new column called 'ENG2WGT' that has the engine to weight ratio

```
[22]: cars['ENG2WGT'] = cars['ENG']/cars['WGT']
```

```
[23]: cars
```

```
[23]:
```

	MPG	CYL	ENG	WGT	ENG2WGT
0	18.0	8	307.0	3504	0.087614
1	15.0	8	350.0	3693	0.094774
2	18.0	8	318.0	3436	0.092549
3	16.0	8	304.0	3433	0.088552
4	17.0	8	302.0	3449	0.087562
..	...	...	...	...	...
387	27.0	4	140.0	2790	0.050179
388	44.0	4	97.0	2130	0.045540
389	32.0	4	135.0	2295	0.058824
390	28.0	4	120.0	2625	0.045714
391	31.0	4	119.0	2720	0.043750

[392 rows x 5 columns]

0.1.4 Question 4: For the kaggle uber other federal.csv data set:

```
[24]: uber = pd.read_csv('kaggle-uber-other-federal.csv',  
                        parse_dates = ['Date', 'Time'])  
uber.head(2)
```

```
[24]:
```

	Date	Time	PU_Address \
0	2014-07-01	2021-07-16 07:15:00	Brooklyn Museum, 200 Eastern Pkwy., BK NY;
1	2014-07-01	2021-07-16 07:30:00	33 Robert Dr., Short Hills NJ;

	DO_Address \
0	1 Brookdale Plaza, BK NY;
1	John F Kennedy International Airport, vitona A...

	Routing Details \
0	PU: Brooklyn Museum, 200 Eastern Pkwy., BK NY;...
1	PU: 33 Robert Dr., Short Hills NJ; DO: John F ...

	PU_Address.1	Status
0	Brooklyn Museum, 200 Eastern Pkwy., BK NY; DO:...	Cancelled
1	33 Robert Dr., Short Hills NJ; DO: John F Kenn...	Arrived

(a) Create a new data frame containing 'Time', 'Status', and 'PU\_Address' columns

```
[25]: new_uber = uber[['Time', 'Status', 'PU_Address']]
new_uber
```

```
[25]:
```

	Time	Status	\
0	2021-07-16 07:15:00	Cancelled	
1	2021-07-16 07:30:00	Arrived	
2	2021-07-16 08:00:00	Assigned	
3	2021-07-16 09:00:00	Assigned	
4	2021-07-16 09:30:00	Assigned	
..	...	...	
94	2021-07-16 06:00:00	Assigned	
95	2021-07-16 08:30:00	Cancelled	
96	2021-07-16 12:00:00	Arrived	
97	2021-07-16 16:45:00	Assigned	
98	2021-07-16 13:30:00	Arrived	

	PU_Address
0	Brooklyn Museum, 200 Eastern Pkwy., BK NY;
1	33 Robert Dr., Short Hills NJ;
2	60 Glenmore Ave., BK NY;
3	128 East 31 St., BK NY;
4	139-39 35 Ave., Flushing NY;
..	...
94	266 prospect park west, brooklyn NY;
95	42 St., BK NY;
96	663 51st Street, BK NY;
97	255 Fieldston Terrace, Bronx NY;
98	Columbia University, 630 W 168 St., NY NY; ST:...

[99 rows x 3 columns]

**(b) Create a new column 'Hour' extracting hour information from 'Time'. Hint: make sure 'Time' is the correct data type and keep the date added.**

```
[26]: df_uber = new_uber.copy()
df_uber
```

```
[26]:
```

	Time	Status	\
0	2021-07-16 07:15:00	Cancelled	
1	2021-07-16 07:30:00	Arrived	
2	2021-07-16 08:00:00	Assigned	
3	2021-07-16 09:00:00	Assigned	
4	2021-07-16 09:30:00	Assigned	
..	...	...	
94	2021-07-16 06:00:00	Assigned	
95	2021-07-16 08:30:00	Cancelled	
96	2021-07-16 12:00:00	Arrived	
97	2021-07-16 16:45:00	Assigned	

98 2021-07-16 13:30:00 Arrived

```

                                PU_Address
0      Brooklyn Museum, 200 Eastern Pkwy., BK NY;
1              33 Robert Dr., Short Hills NJ;
2              60 Glenmore Ave., BK NY;
3              128 East 31 St., BK NY;
4              139-39 35 Ave., Flushing NY;
..
94      266 prospect park west, brooklyn NY;
95              42 St., BK NY;
96              663 51st Street, BK NY;
97      255 Fieldston Terrace, Bronx NY;
98 Columbia University, 630 W 168 St., NY NY; ST:...
```

[99 rows x 3 columns]

```
[27]: df_uber['Hour'] = df_uber['Time'].dt.hour
df_uber
```

```
[27]:
      Time      Status \
0  2021-07-16 07:15:00 Cancelled
1  2021-07-16 07:30:00 Arrived
2  2021-07-16 08:00:00 Assigned
3  2021-07-16 09:00:00 Assigned
4  2021-07-16 09:30:00 Assigned
..
94 2021-07-16 06:00:00 Assigned
95 2021-07-16 08:30:00 Cancelled
96 2021-07-16 12:00:00 Arrived
97 2021-07-16 16:45:00 Assigned
98 2021-07-16 13:30:00 Arrived
```

```

                                PU_Address  Hour
0      Brooklyn Museum, 200 Eastern Pkwy., BK NY;      7
1              33 Robert Dr., Short Hills NJ;      7
2              60 Glenmore Ave., BK NY;      8
3              128 East 31 St., BK NY;      9
4              139-39 35 Ave., Flushing NY;      9
..
94      266 prospect park west, brooklyn NY;      6
95              42 St., BK NY;      8
96              663 51st Street, BK NY;     12
97      255 Fieldston Terrace, Bronx NY;     16
98 Columbia University, 630 W 168 St., NY NY; ST:...    13
```

[99 rows x 4 columns]



**(c) Set the index of the data frame to 'Time'.**

```
[28]: df_uber.set_index('Time', inplace = True)
df_uber
```

```
[28]:
```

	Status \	
Time		
2021-07-16 07:15:00	Cancelled	
2021-07-16 07:30:00	Arrived	
2021-07-16 08:00:00	Assigned	
2021-07-16 09:00:00	Assigned	
2021-07-16 09:30:00	Assigned	
...	...	
2021-07-16 06:00:00	Assigned	
2021-07-16 08:30:00	Cancelled	
2021-07-16 12:00:00	Arrived	
2021-07-16 16:45:00	Assigned	
2021-07-16 13:30:00	Arrived	

	PU_Address	Hour
Time		
2021-07-16 07:15:00	Brooklyn Museum, 200 Eastern Pkwy., BK NY;	7
2021-07-16 07:30:00	33 Robert Dr., Short Hills NJ;	7
2021-07-16 08:00:00	60 Glenmore Ave., BK NY;	8
2021-07-16 09:00:00	128 East 31 St., BK NY;	9
2021-07-16 09:30:00	139-39 35 Ave., Flushing NY;	9
...	...	...
2021-07-16 06:00:00	266 prospect park west, brooklyn NY;	6
2021-07-16 08:30:00	42 St., BK NY;	8
2021-07-16 12:00:00	663 51st Street, BK NY;	12
2021-07-16 16:45:00	255 Fieldston Terrace, Bronx NY;	16
2021-07-16 13:30:00	Columbia University, 630 W 168 St., NY NY; ST:...	13

[99 rows x 3 columns]

**(d) Display records at positions between 10 and 15 inclusively**

```
[29]: df_uber.iloc[10:16]
```

```
[29]:
```

	Status	PU_Address	Hour
Time			
2021-07-16 20:00:00	Assigned	35-36 32 St., Astoria NY;	20
2021-07-16 03:30:00	Arrived	862 East 21 Street, BK NY;	3
2021-07-16 14:00:00	Assigned	1539 71st Street, BK NY;	14
2021-07-16 15:00:00	Arrived	208 Elmwood ave, BK NY;	15
2021-07-16 20:45:00	Arrived	543 1 St., BK NY;	20
2021-07-16 05:00:00	Arrived	513 Montgomery StreetBK NY;	5

(e) Display the 'PU\_Address' for records whose index is 'Today's date 20:00:00': hint: today's date depends on the day you are working on the question.

```
[30]: df_uber.index.hour == 7
```

```
[30]: array([ True,  True, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False, False,
         True, False, False, False, False, False, False, False, False, False,
        False, False, False, False,  True, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False, False,
         True,  True, False, False, False, False, False, False,  True, False,
        False, False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False, False])
```

```
[31]: df_uber[(df_uber.index.hour == 20) & (df_uber.index.minute == 0)]
```

```
[31]:
```

	Status	PU_Address	Hour
Time			
2021-07-16 20:00:00	Assigned	35-36 32 St., Astoria NY;	20
2021-07-16 20:00:00	Arrived	717 President Street, BK NY;	20

(f) Reset the data frame so that 'Time' is a column again

```
[32]: df_uber.reset_index('Time', inplace = True)
```

```
[33]: df_uber
```

```
[33]:
```

	Time	Status	\
0	2021-07-16 07:15:00	Cancelled	
1	2021-07-16 07:30:00	Arrived	
2	2021-07-16 08:00:00	Assigned	
3	2021-07-16 09:00:00	Assigned	
4	2021-07-16 09:30:00	Assigned	
..	...	...	
94	2021-07-16 06:00:00	Assigned	
95	2021-07-16 08:30:00	Cancelled	
96	2021-07-16 12:00:00	Arrived	
97	2021-07-16 16:45:00	Assigned	
98	2021-07-16 13:30:00	Arrived	

	PU_Address	Hour
0	Brooklyn Museum, 200 Eastern Pkwy., BK NY;	7
1	33 Robert Dr., Short Hills NJ;	7
2	60 Glenmore Ave., BK NY;	8
3	128 East 31 St., BK NY;	9
4	139-39 35 Ave., Flushing NY;	9
..	...	...

```

94          266 prospect park west, brooklyn NY;      6
95                                42 St., BK NY;      8
96                                663 51st Street, BK NY; 12
97          255 Fieldston Terrace, Bronx NY;      16
98 Columbia University, 630 W 168 St., NY NY; ST:... 13

```

```
[99 rows x 4 columns]
```

**0.1.5 Question 5:** You work for a large school district as a data analyst. Your boss wants to purchase a large amount of cereal for school breakfasts. He needs to choose a manufacturer and product. He wants you to prepare a presentation for the executive team. You are provided with some data in `cereal.csv`:

(a) Initially explore the dataset by looking at the number of records, column names, column types and few record values. Through this initial look, combined with your boss's goal above, list five analytics queries or questions that you would have about this dataset in your exploratory process.

```
[34]: master_cereal = pd.read_csv('cereal.csv')
      master_cereal.head()
```

```
[34]:
```

		name	mfr	type	calories	protein	fat	sodium	fiber	\
0		100% Bran	N	C	70	4	1	130	10.0	
1		100% Natural Bran	Q	C	120	3	5	15	2.0	
2		All-Bran	K	C	70	4	1	260	9.0	
3	All-Bran with Extra Fiber		K	C	50	4	0	140	14.0	
4		Almond Delight	R	C	110	2	2	200	1.0	

	carbo	sugars	potass	vitamins	shelf	weight	cups	rating
0	5.0	6	280	25	3	1.0	0.33	68.402973
1	8.0	8	135	0	3	1.0	1.00	33.983679
2	7.0	5	320	25	3	1.0	0.33	59.425505
3	8.0	0	330	25	3	1.0	0.50	93.704912
4	14.0	8	-1	25	3	1.0	0.75	34.384843

```
[35]: master_cereal.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 77 entries, 0 to 76
Data columns (total 16 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   name        77 non-null    object
 1   mfr         77 non-null    object
 2   type        77 non-null    object
 3   calories    77 non-null    int64
 4   protein     77 non-null    int64
 5   fat         77 non-null    int64

```

```

6   sodium      77 non-null    int64
7   fiber       77 non-null    float64
8   carbo       77 non-null    float64
9   sugars      77 non-null    int64
10  potass      77 non-null    int64
11  vitamins    77 non-null    int64
12  shelf       77 non-null    int64
13  weight      77 non-null    float64
14  cups        77 non-null    float64
15  rating      77 non-null    float64
dtypes: float64(5), int64(8), object(3)
memory usage: 9.8+ KB

```

**Questions** 1. Which manufacturers produce the top five highest rating cereals? 2. Which manufacturer produces healthier cereal such that it is high in vitamins and low in sugar? 3. Which type of cereal has a higher rating in average? 4. What is the average cups in a serving for a typical cereal? 5. Would more cups in a serving be more popular?

**(b) To answer the questions you have listed above, what columns from this data you would need? Print some basic statistical summaries and plots for at least 5 columns you would need to check missing values, extreme values and value distributions. List three findings you have at this point.** > In order to answer the questions I have listed above, I would need the columns *mfr*, *type*, *sugars*, *vitamins*, *cups*, and *rating*. > - *mfr*: Manufacturer of cereal: - A = American Home Food Products - G = General Mills - K = Kelloggs - N = Nabisco - P = Post - Q = Quaker Oats - R = Ralston Purina > - *type*: cold/hot > - *sugars*: grams of sugars > - *vitamins*: vitamins and minerals - 0, 25, or 100, indicating the typical percentage of FDA recommended > - *cups*: number of cups in one serving > - *rating*: a rating of the cereals

```
[36]: master_cereal.columns
```

```
[36]: Index(['name', 'mfr', 'type', 'calories', 'protein', 'fat', 'sodium', 'fiber',
        'carbo', 'sugars', 'potass', 'vitamins', 'shelf', 'weight', 'cups',
        'rating'],
        dtype='object')
```

```
[37]: df_cereal = master_cereal[['mfr', 'type', 'sugars', 'vitamins', 'cups',
        → 'rating']]
```

```
[38]: df_cereal.head()
```

```
[38]:   mfr type  sugars  vitamins  cups  rating
0    N   C        6         25  0.33  68.402973
1    Q   C        8          0  1.00  33.983679
2    K   C        5         25  0.33  59.425505
3    K   C        0         25  0.50  93.704912
4    R   C        8         25  0.75  34.384843
```

```
[39]: df_cereal.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 77 entries, 0 to 76
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  -
0   mfr          77 non-null    object
1   type         77 non-null    object
2   sugars       77 non-null    int64
3   vitamins     77 non-null    int64
4   cups         77 non-null    float64
5   rating       77 non-null    float64
dtypes: float64(2), int64(2), object(2)
memory usage: 3.7+ KB

```

```
[40]: df_cereal.describe()
```

```

[40]:      sugars  vitamins  cups  rating
count  77.000000  77.000000  77.000000  77.000000
mean     6.922078   28.246753   0.821039  42.665705
std     4.444885   22.342523   0.232716  14.047289
min    -1.000000    0.000000   0.250000  18.042851
25%     3.000000   25.000000   0.670000  33.174094
50%     7.000000   25.000000   0.750000  40.400208
75%    11.000000   25.000000   1.000000  50.828392
max    15.000000  100.000000   1.500000  93.704912

```

```
[41]: df_cereal['sugars'].value_counts().loc[-1]
```

```
[41]: 1
```

```
[42]: df_cereal['sugars'].value_counts()
```

```


[42]: 3      13
      0      7
      12     7
      6      7
      10     5
      5      5
      8      5
      11     5
      9      4
      7      4
      13     4
      14     3
      2      3
      15     2
      4      1
      1      1
     -1      1
Name: sugars, dtype: int64

```

```
[43]: fig, ax = plt.subplots(figsize = (6,4))
```

```
ax.boxplot(df_cereal['sugars'])
```

```
[43]: {'whiskers': [<matplotlib.lines.Line2D at 0x7fe799e14d30>,  
                 <matplotlib.lines.Line2D at 0x7fe799e14e10>],  
      'caps': [<matplotlib.lines.Line2D at 0x7fe799e2d400>,  
              <matplotlib.lines.Line2D at 0x7fe799e2d748>],  
      'boxes': [<matplotlib.lines.Line2D at 0x7fe799e14940>],  
      'medians': [<matplotlib.lines.Line2D at 0x7fe799e2da90>],  
      'fliers': [<matplotlib.lines.Line2D at 0x7fe799e2ddd8>],  
      'means': []}
```



**Observations** - In the *sugars* column, the minimum grams of sugar is -1 which is not logical as the least amount of sugar anything could contain is 0 gram. After delving more into this weird phenonme, I noticed there was only one element of -1 out of 77 elements. On the boxplot of `df_cereal['sugars']`, the one element of -1 didn't caused a big impact. It wasn't necessary to replace the element, however, it would be better to change it from -1 to 0 so the grams of sugars would make more logical sense. - The *mfr* column is currently in object dtype, however, changing the *mfr* column from object type to category dtype would help more in the later data exploration portion. - The *type* column is also currently in object dtype, however, changing the dtype of it from object to categorical dtype would make exploring the type of cereal more efficient. - The *mfr* column contains the abbreviation of the manufacturers' names. Thus, I would need to keep referring back to above state-

ments for clarification on the abbreviation. In order to solve that, I would substitute back the actual names for the abbreviated names.

```
[44]: cereal = df_cereal.copy()
```

```
[45]: cereal[cereal['sugars']== -1].index
```

```
[45]: Int64Index([57], dtype='int64')
```

```
[46]: cereal['sugars'].iloc[57] = 0
```

```
/Users/amandaliluo/anaconda3/lib/python3.7/site-  
packages/pandas/core/indexing.py:671: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
self._setitem_with_indexer(indexer, value)
```

```
[47]: cereal.iloc[57]
```

```
[47]: mfr          Q  
      type       H  
      sugars     0  
      vitamins   0  
      cups       0.67  
      rating    50.8284  
      Name: 57, dtype: object
```

```
[48]: cereal['mfr'] = cereal['mfr'].astype('category')  
      cereal['type'] = cereal['type'].astype('category')
```

```
[49]: cereal.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 77 entries, 0 to 76  
Data columns (total 6 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   mfr          77 non-null      category  
1   type         77 non-null      category  
2   sugars       77 non-null      int64  
3   vitamins     77 non-null      int64  
4   cups         77 non-null      float64  
5   rating       77 non-null      float64  
dtypes: category(2), float64(2), int64(2)  
memory usage: 3.1 KB
```

```
[50]: cereal['type'].unique()
```

```
[50]: [C, H]
Categories (2, object): [C, H]
```

```
[51]: cereal['vitamins'].unique()
```

```
[51]: array([ 25,   0, 100])
```

```
[52]: cereal['mfr'].value_counts()
```

```
[52]: K    23
      G    22
      P     9
      R     8
      Q     8
      N     6
      A     1
      Name: mfr, dtype: int64
```

```
[53]: #x = {'A': 'American Home Food Products',
        #'G': 'General Mills', 'K': 'Kelloggs',
        ##Q': 'Quaker Oats', 'R': 'Ralston Purina'}
name = pd.DataFrame({'mfr': ['A', 'G', 'K', 'N', 'P', 'Q', 'R'],
                    'manufacturer': ['American Home Food Products', 'General_
→Mills',
                                   'Kelloggs', 'Nabisco', 'Post', 'Quaker Oats',_
→'Ralston Purina']})

df = pd.merge(cereal, name, on = 'mfr')
df = df.drop('mfr', axis =1)
df = df[['manufacturer', 'type', 'sugars', 'vitamins', 'cups', 'rating']]
df['manufacturer'] = df['manufacturer'].astype('category')
df
```

```
[53]:
```

	manufacturer	type	sugars	vitamins	cups	rating
0	Nabisco	C	6	25	0.33	68.402973
1	Nabisco	H	0	0	1.00	64.533816
2	Nabisco	C	0	0	1.00	68.235885
3	Nabisco	C	0	0	0.67	74.472949
4	Nabisco	C	0	0	0.67	72.801787
...	...	...	...	...	...	...
72	Post	C	3	25	0.25	53.371007
73	Post	C	4	25	0.33	45.811716
74	Post	C	11	25	1.33	28.742414
75	Post	C	14	25	0.67	37.840594
76	American Home Food Products	H	3	25	1.00	54.850917

```
[77 rows x 6 columns]
```



(c) To answer the questions you have listed above, describe 5 visualizations that can help you explore this data. (Example: line graph between variable x and y).

(d) Plot at three DIFFERENT graphs you have mentioned above. Summarise some findings from each plot. (For example, line chart of x and y and line chart of a and b doesn't count as different; line chart and box plot does. )

1. Which manufacturers produces the top five highest rating cereals? - Based on the histogram and the table below, Kelloggs produced the highest rating cereal of 93.70 and Nabisco produced the other 4 high rating cereals of 74.47, 72.80, 72.80, 68.40, and 68.24. Even though Kelloggs produced the highest rating cereal, however, Nabisco's productions composed majority of the top five highest rating cereals.

```
[54]: x = df['rating'].sort_values(ascending = False).head(5).index
x
```

```
[54]: Int64Index([15, 3, 4, 0, 2], dtype='int64')
```

```
[55]: top5 = df.iloc[x]
top5
```

```
[55]:
```

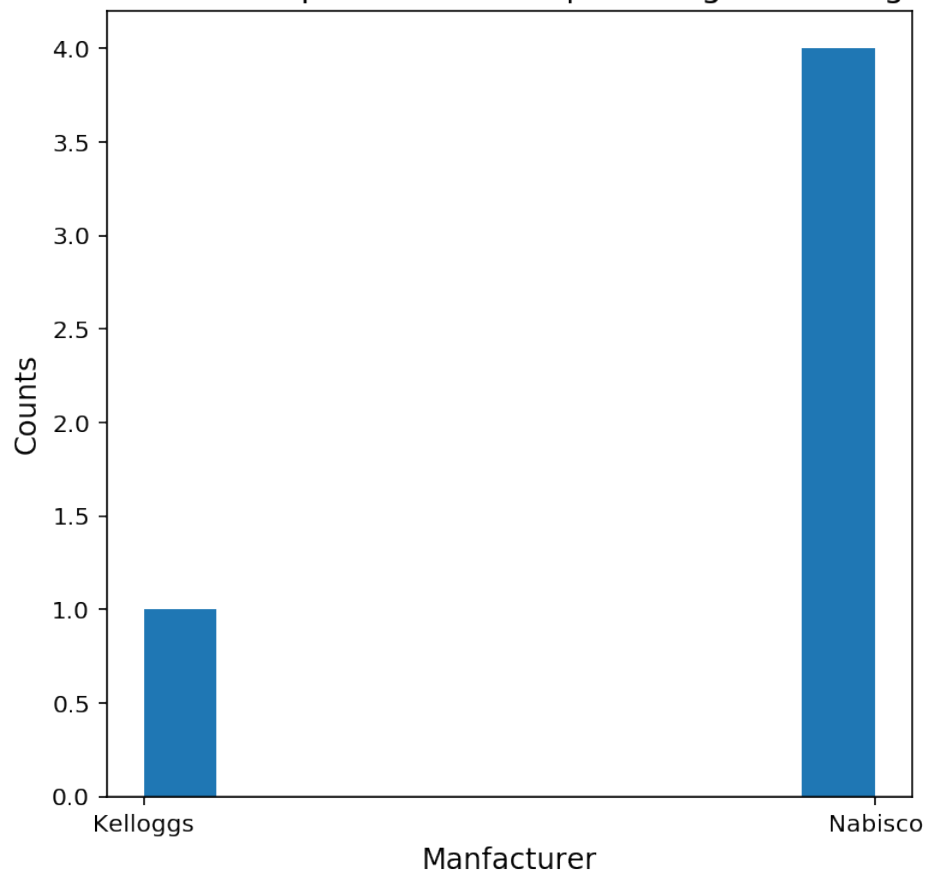
	manufacturer	type	sugars	vitamins	cups	rating
15	Kelloggs	C	0	25	0.50	93.704912
3	Nabisco	C	0	0	0.67	74.472949
4	Nabisco	C	0	0	0.67	72.801787
0	Nabisco	C	6	25	0.33	68.402973
2	Nabisco	C	0	0	1.00	68.235885

```
[56]: fig, ax = plt.subplots(figsize = (6,6))

ax.hist(top5.manufacturer)
ax.set_xlabel('Manufacturer', fontsize = 12)
ax.set_ylabel('Counts', fontsize = 12)
ax.set_title('Which manufacturers produces the top five highest rating cereals?',
             →fontsize = 14)

plt.show()
```

Which manufacturers produce the top five highest rating cereals?



**2. Which manufacturer produces healthier cereal that is high in vitamins and low in sugar?** - Based on the graph below and the table, American Home Food Products produces the healthiest cereal in average compared to other manufacturers. It has the highest ratio of vitamins to sugars of around 8.33. While Kellogg's, who has the highest rating, claimed the third place with a ratio of around 6.79.

```
[57]: df1 = df.copy()
df1['vitamins2sugars'] = df1['vitamins']/df1['sugars']
df1['vitamins2sugars'].unique()
# since nan means we have 0 in vitamin and 0 in grams of sugar, and inf means 0
  ↳ in sugar
# I will replace nan and inf with zero
```

```
[57]: array([ 4.16666667,      nan,  5.         ,  0.         ,  2.08333333,
          2.27272727,      inf,  1.78571429, 12.5         ,  3.57142857,
          8.33333333,  1.92307692, 16.66666667, 11.11111111,  2.77777778,
          33.33333333,  1.66666667,  3.125         ,  2.5         , 25.         ,
          7.14285714,  6.25         ])
```

```
[58]: df1[df1['vitamins2sugars'].isnull()]
```

```
[58]:
```

	manufacturer	type	sugars	vitamins	cups	rating	vitamins2sugars
1	Nabisco	H	0	0	1.00	64.533816	NaN
2	Nabisco	C	0	0	1.00	68.235885	NaN
3	Nabisco	C	0	0	0.67	74.472949	NaN
4	Nabisco	C	0	0	0.67	72.801787	NaN
10	Quaker Oats	C	0	0	1.00	60.756112	NaN
11	Quaker Oats	C	0	0	1.00	63.005645	NaN
13	Quaker Oats	H	0	0	0.67	50.828392	NaN

```
[59]: df1['vitamins2sugars'].sort_values(ascending = False)
```

```
[59]:
```

15	inf
60	33.333333
31	33.333333
62	33.333333
47	25.000000
...	
3	NaN
4	NaN
10	NaN
11	NaN
13	NaN

Name: vitamins2sugars, Length: 77, dtype: float64

```
[60]: df1.iloc[15]
```

```
[60]:
```

manufacturer	Kelloggs
type	C
sugars	0
vitamins	25
cups	0.5
rating	93.7049
vitamins2sugars	inf

Name: 15, dtype: object

```
[61]: df1 = df1.drop(index =15)  
df1 = df1.dropna()
```

```
[62]: df1
```

```
[62]:
```

	manufacturer	type	sugars	vitamins	cups	rating	\
0	Nabisco	C	6	25	0.33	68.402973	
5	Nabisco	C	5	25	1.00	59.363993	
6	Quaker Oats	C	8	0	1.00	33.983679	
7	Quaker Oats	C	12	25	0.75	18.042851	
8	Quaker Oats	C	11	25	1.00	21.871292	
..	...	...	...	...	...	...	
72	Post	C	3	25	0.25	53.371007	
73	Post	C	4	25	0.33	45.811716	

74		Post	C	11	25	1.33	28.742414
75		Post	C	14	25	0.67	37.840594
76	American Home Food Products		H	3	25	1.00	54.850917

	vitamins	sugars
0	4.166667	
5	5.000000	
6	0.000000	
7	2.083333	
8	2.272727	
..	...	
72	8.333333	
73	6.250000	
74	2.272727	
75	1.785714	
76	8.333333	

[69 rows x 7 columns]

```
[63]: barh = df1.groupby('manufacturer').mean()
barh
```

```
[63]:
```

	sugars	vitamins	cups	rating \
manufacturer				
American Home Food Products	3.000000	25.000000	1.000000	54.850917
General Mills	7.954545	35.227273	0.875000	34.485852
Kelloggs	7.909091	35.227273	0.809545	41.780896
Nabisco	5.500000	25.000000	0.665000	63.883483
Post	8.777778	25.000000	0.714444	41.705744
Quaker Oats	8.600000	20.000000	0.784000	33.747554
Ralston Purina	6.125000	25.000000	0.871250	41.542997

	vitamins	sugars
manufacturer		
American Home Food Products	8.333333	
General Mills	7.476863	
Kelloggs	6.799596	
Nabisco	4.583333	
Post	3.876864	
Quaker Oats	2.537879	
Ralston Purina	5.750473	

```
[64]: barh.index
```

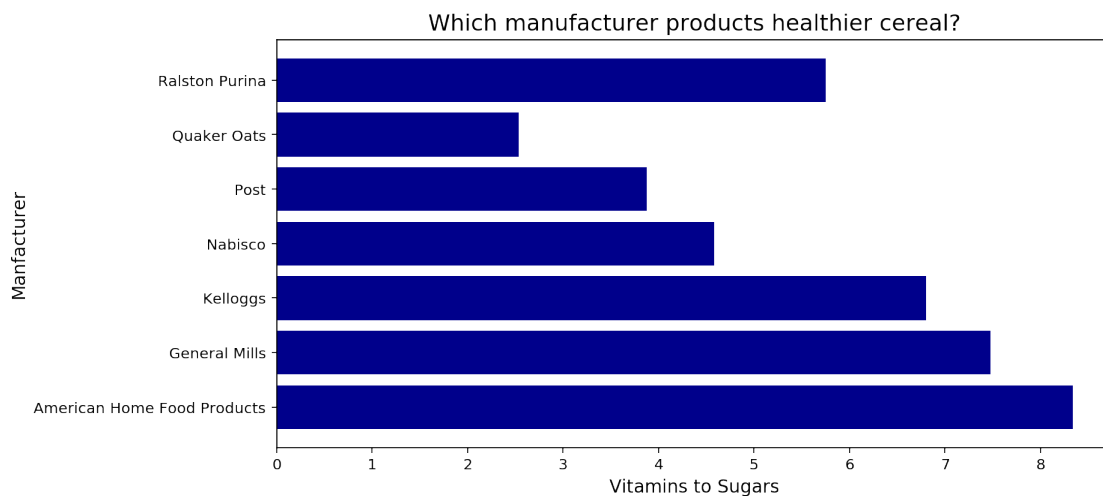
```
[64]: CategoricalIndex(['American Home Food Products', 'General Mills', 'Kelloggs',
                        'Nabisco', 'Post', 'Quaker Oats', 'Ralston Purina'],
                        categories=['American Home Food Products', 'General Mills',
                        'Kelloggs', 'Nabisco', 'Post', 'Quaker Oats', 'Ralston Purina'], ordered=False,
                        name='manufacturer', dtype='category')
```

```
[65]: fig, ax = plt.subplots(figsize =(10,5))

x = barh.index
y = barh.vitamins2sugars

ax.barh(x, y, color= 'darkblue')
ax.set_ylabel('Manufacturer', fontsize = 12)
ax.set_xlabel('Vitamins to Sugars', fontsize = 12)
ax.set_title('Which manufacturer products healthier cereal?', fontsize = 15)

plt.show()
```



**3. Which type of cereal has a higher rating in average?** - Hot cereals have a higher rating of around 56.74 in average compared to cold cereals.

```
[66]: barh2 = df.groupby('type').mean()
barh2
```

```
[66]:
```

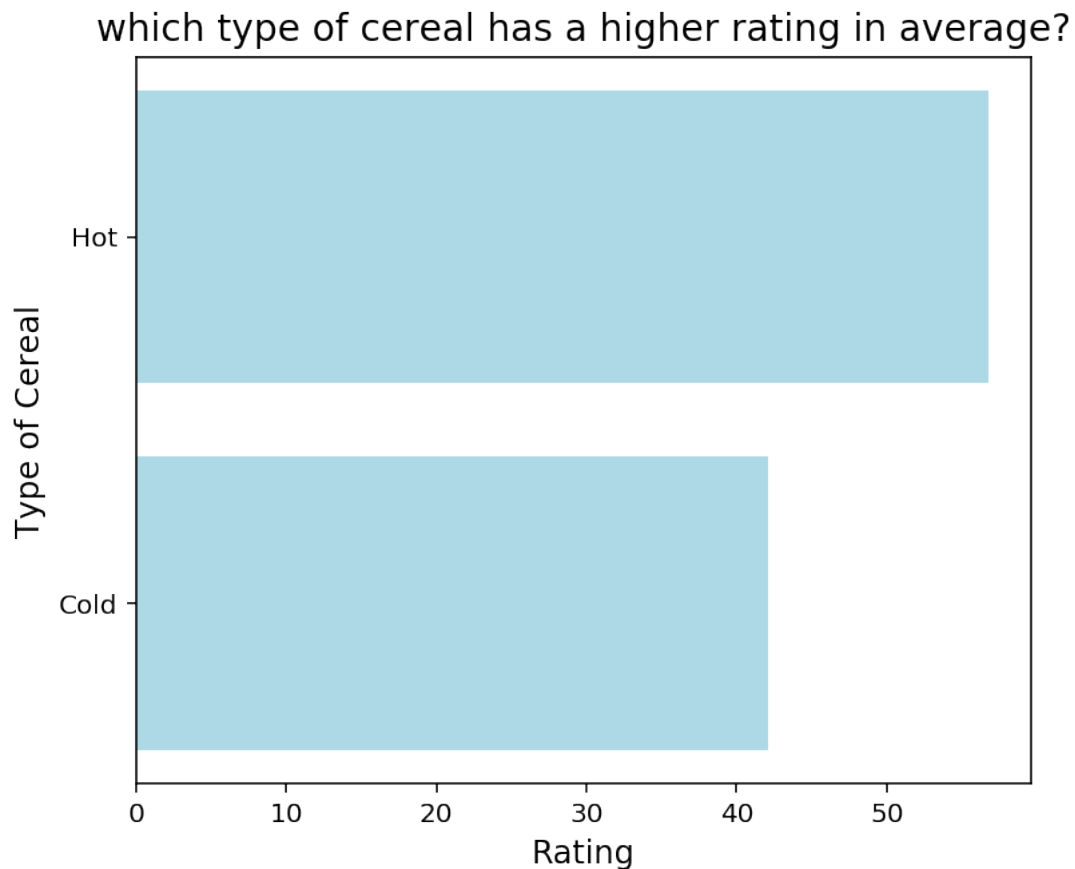
	sugars	vitamins	cups	rating
type				
C	7.175676	29.054054	0.818243	42.095218
H	1.000000	8.333333	0.890000	56.737708

```
[67]: fig, ax = plt.subplots(figsize = (6,5))

x = barh2.index
y = barh2.rating

ax.barh(x, y, color = 'lightblue')
ax.set_xlabel('Rating', fontsize = 12)
ax.set_ylabel('Type of Cereal', fontsize = 12)
```

```
ax.set_title('which type of cereal has a higher rating in average?', fontsize = 14)
ax.set_yticklabels(['Cold', 'Hot'])
plt.show()
```



**4. What is the average cups in a serving for a typical cereal?** - The typical cereal have an average of around 0.82 cups of serving.

```
[68]: df['cups'].mean()
```

```
[68]: 0.8210389610389613
```

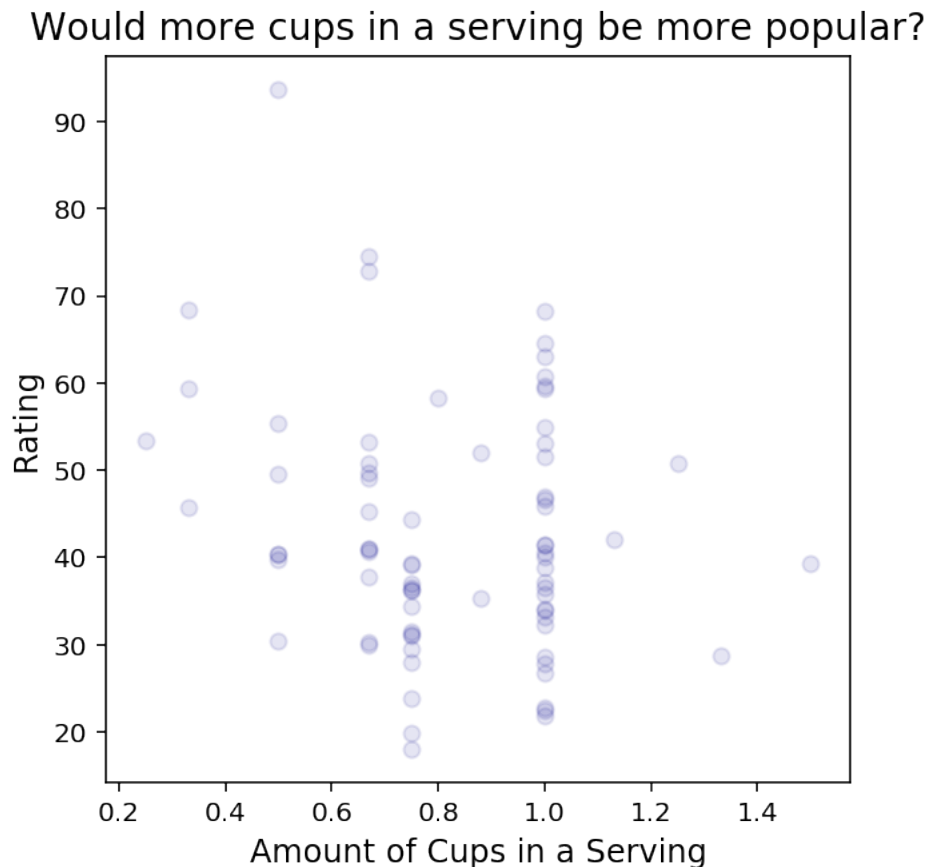
5. Would more cups in a serving be more popular?

- From the scatterplot below, cereals that have more cups in a serving have lower ratings. However, the amount of cups in a serving and the rating of the cereal has a weak neagative correlation of around -0.20. Thus, the amount of cups in a serving have little effect on the rating of the cereal.

```
[69]: f = df[['cups', 'rating']].corr()  
f
```

```
[69]:      cups  rating  
cups    1.00000 -0.20316  
rating -0.20316  1.00000
```

```
[70]: fig, ax = plt.subplots(figsize=(5,5))  
x = df['cups']  
y = df['rating']  
  
ax.scatter(x, y, color = 'darkblue', alpha = 0.1)  
ax.set_xlabel('Amount of Cups in a Serving', fontsize = 12)  
ax.set_ylabel('Rating', fontsize = 12)  
ax.set_title('Would more cups in a serving be more popular?', fontsize = 14)  
  
plt.show()
```



(e) Besides the information provided in the data, what else information might be helpful to achieve your boss' goal? - Aside from the information in the data set, some

additional information that would help my boss pick a manufacturer and product for school breakfasts. One additional piece of information that would be helpful is if the cereal is being used for school breakfasts before. Another one is if the cereal has received complaints before in amounts.